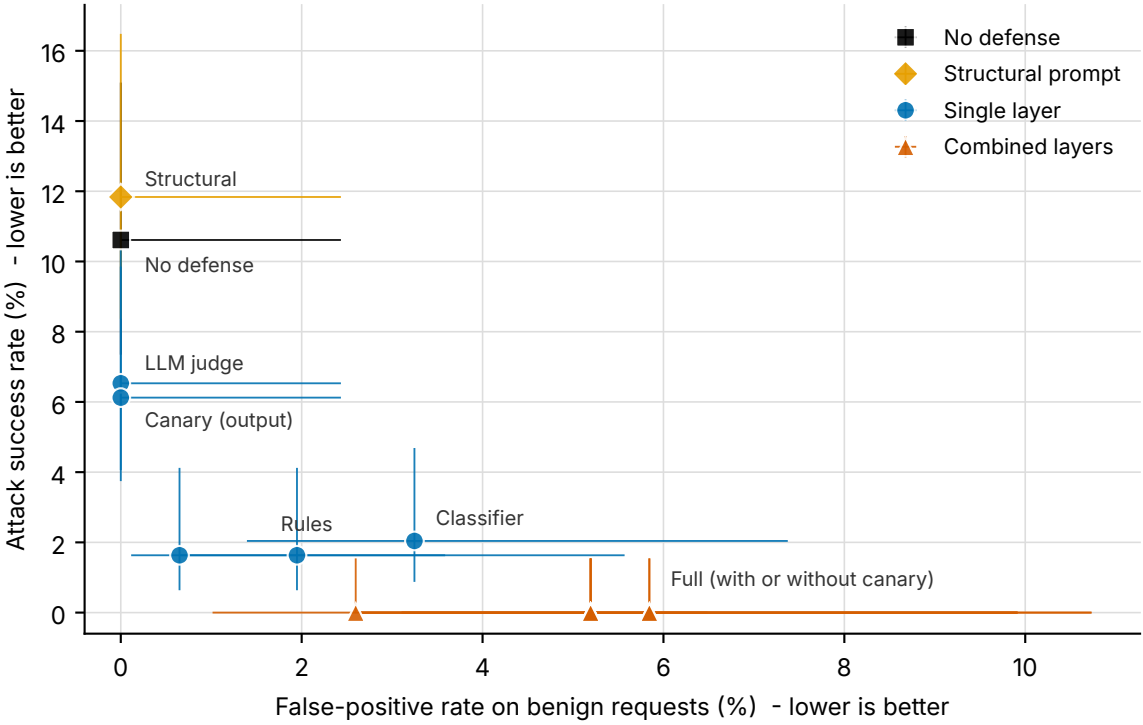
No defense
Structural prompt
Single layer
Combined layers
Structural
No defense
LLM judge
Canary (output)
Rules
Classifier
Full (with or without canary)
Attack success rate (%) - lower is better
False-positive rate on benign requests (%) - lower is better
0
2
4
6
8
10
12
14
16