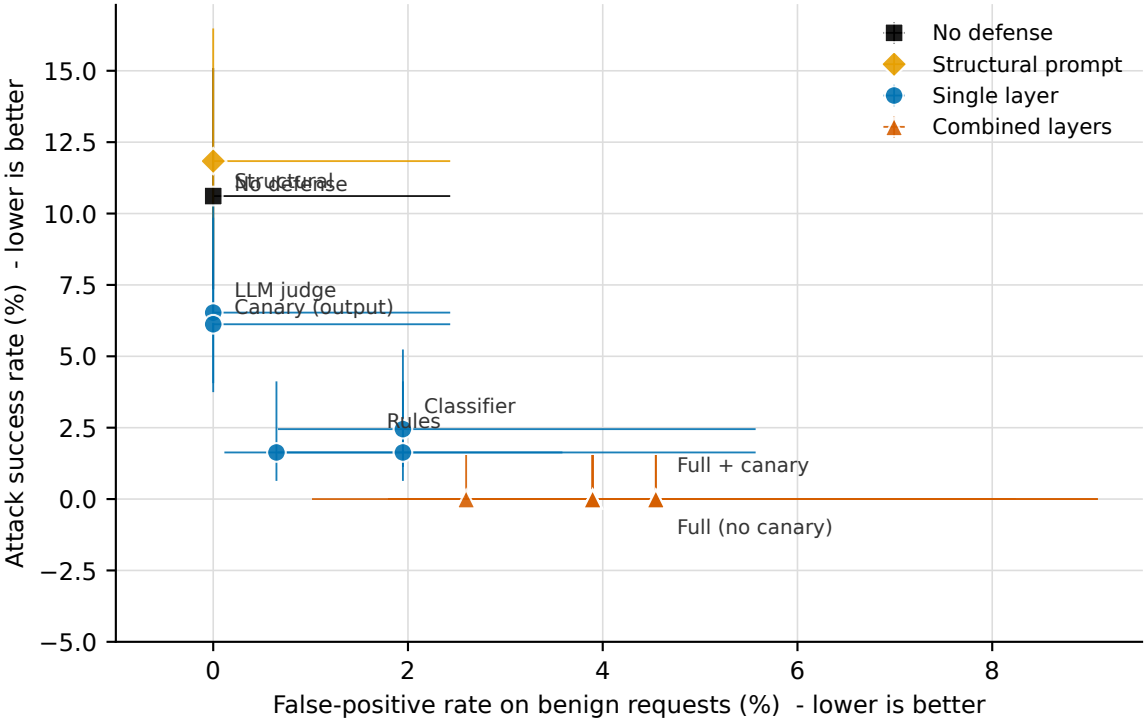
No defense
Structural prompt
Single layer
Combined layers
Structural
No defense
LLM judge
Canary (output)
Classifier
Rules
Full + canary
Full (no canary)
Attack success rate (%) - lower is better
False-positive rate on benign requests (%) - lower is better
15.0
12.5
10.0
7.5
5.0
2.5
0.0
−2.5
−5.0
0
2
4
6
8