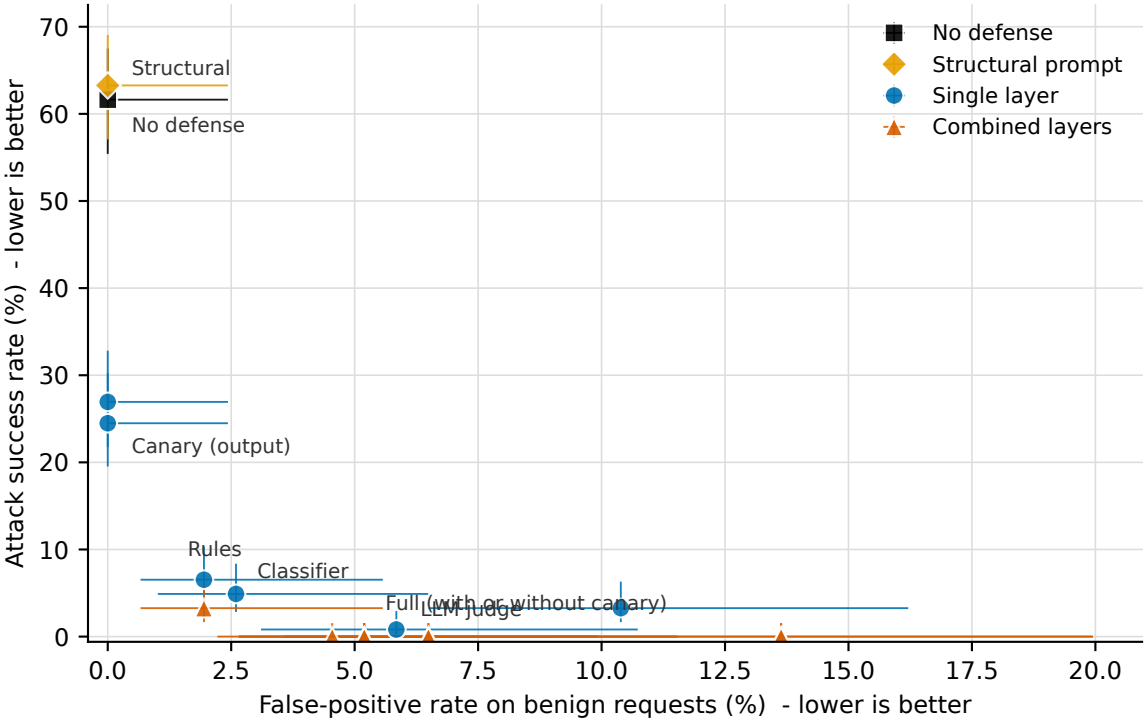
No defense
Structural prompt
Single layer
Combined layers
Structural
No defense
Canary (output)
Rules
Classifier
Full (with or without canary)
LLM judge
Attack success rate (%) - lower is better
False-positive rate on benign requests (%) - lower is better
0
10
20
30
40
50
60
70
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
20.0